

hvsat slides

SAT

Boolean Satisfiability

Boolean formulas

- Literals / Variables: x_1, x_2, x_3
- Negation: $\neg x_1$
- Conjunction: $x_1 \wedge x_2$
- Disjunction: $x_1 \vee x_2$

Boolean formulas

- CNF: conjunction of disjunctions
 - $(x_1 \vee x_2 \vee \neg x_3) \wedge (\neg x_2 \vee x_3)$
- DNF: disjunction of conjunctions
 - $(x_2 \wedge x_3) \vee (x_1 \wedge \neg x_2) \vee (\neg x_2 \wedge \neg x_3)$
- “DNF is easier to think about, but tends to blow up in space”
- “CNF is cheap to construct, but difficult to reason about”

Boolean Satisfiability Problem (SAT)

- Input: boolean formula, usually in CNF
 - Example: $(x_1 \vee x_2 \vee \neg x_3) \wedge (\neg x_2 \vee x_3)$
 - *Literals & Clauses*
- Output: is there a satisfying assignment to all variables that makes the formula hold?
- If there is, return SATISFIABLE and the assignments.
- If there is not, return UNSATISFIABLE.
 - Example: Satisfiable! For instance: $x_1 = \text{False}, x_2 = \text{True}, x_3 = \text{True}$
 - $(\textcolor{red}{x}_1 \vee \textcolor{green}{x}_2 \vee \neg \textcolor{red}{x}_3) \wedge (\neg \textcolor{red}{x}_2 \vee \textcolor{green}{x}_3)$

First implementation: Loops

```
for x1 in [True, False]:
```

```
    for x2 in [True, False]:
```

```
        if formula(x1, x2) return SAT
```

```
return UNSAT
```

- Will we return SAT if there is a solution? Yes!
- Will we return UNSAT if there is no solution? Yes!

Implementation: Loops.py


- Avert thine eyes

Why is SAT interesting? (1 / 3)

- It was the first problem proven to be NP-complete.
 - NP-complete: it scales worse than polynomial time
- With n variables, the search space is 2^n
- Other NP-complete problems:
 - Knapsack
 - Hamiltonian Path
 - ...
- *SAT is the simplest hard problem*

SAT is a
baby
problem!


It looks cute,
but is hard
to handle!



Why is SAT interesting? (2 / 3)

- SAT is intractable in theory
- But in practice, state-of-the-art SAT solvers can solve huge problem instances in almost no time at all!
- SAT competition: solve as many instances as possible in a given time
- Solvers get faster and faster, due to clever engineering and mathematics!

SAT Competition All Time Winners on SAT Competition 2022 Benchmarks



Why is SAT interesting? (3 / 3)

- The SAT enlightenment: Bounded Model Checking (1999)
- Formal verification technique used mainly in EDA community
- BDDs were used before
 - Grow exponentially in complexity
 - Very dependent on input form
- Previously intractable problems were made possible by SAT
- Important solvers: GRASP and CHAFF

Lingua Franca: DIMACS CNF

- File format for CNF formulas.
- All SAT solvers can parse these.
 - Incredibly easy to parse, actually.
- Comments
- Header (nr variables, nr clauses)
- Clauses (1 means x_1 , -2 means $\neg x_2$)

```
c · simple · CNF
c · Satisfiable.
p · cnf · 3 · 2
1 · 2 · -3 · 0
-2 · 3 · 0
```

Same problem in different formats

$$(x_1 \vee x_2 \vee \neg x_3) \wedge (\neg x_2 \vee x_3)$$

Solution:

Satisfiable, model: $[x_1 = F, x_2 = T, x_3 = T]$

$$(x_1 \vee x_2 \vee \neg x_3) \wedge (\neg x_2 \vee x_3)$$

```
c · simple · CNF
c · Satisfiable.
p · cnf · 3 · 2
1 · 2 · -3 · 0
-2 · 3 · 0
```

```
SATISFIABLE
Final model: [-1, 2, 3]
[ x1, x2, -x3]
[-x2, x3]
```

Same problem in different formats

$$x1 \wedge \neg x1$$

Solution:

Unsatisfiable.

```
c · simpler · CNF
c · Unsatisfiable.
p · cnf · 2 · 1
1 · 0
-1 · 0
```

```
UNSATISFIABLE
Final model: [1]
[ x1]
[ -x1]
```

Local search

- AKA “just start flipping bits”
- Assign all variables
- Loop:
 - If SAT return SAT
 - Else negate some assignment

```
for literal in self.literals():  
    assignment = self.random_assignment(literal)  
    self.assign(assignment)  
    for _ in range(timeout):  
        if self.formula_satisfied():  
            return True  
        else:  
            self.flip_random_assignment()  
    else:  
        logging.critical("naive_local_search timeout")  
    return None
```

- Will we return SAT if there is a solution? Yes!
- Will we return UNSAT if there is no solution? No!

Stolen slide 1

1992 - Planning As Satisfiability

$$PAS(S, I, T, G, k) = I(s_0) \wedge \bigwedge_{i=0}^{k-1} T(s_i, s_{i+1}) \wedge \bigvee_{i=0}^k G(s_i)$$

where:

S the set of possible states s_i

I the initial state

T transitions between states

G goal state

k bound

If the formula is **satisfiable**, then there is a **plan** of length at most k .

Stolen slide 2

- The decision procedure is **very simple to implement and very fast!**
- Efficiency depends on which literal to flip, and the values of the parameters.
- Problem with local minima: use of Random Walks!
- Main drawback: **incomplete, cannot answer UNSAT!**
- **Lesson: An agile (fast) SAT solver sometimes better than a clever one!**

Implementation: NaiveLocalSearch.py

- Assigns everything completely at random
- No effort to, for example, target falsified clauses

```
for literal in self.literals():
    assignment = self.random_assignment(literal)
    self.assign(assignment)
for _ in range(timeout):
    if self.formula_satisfied():
        return True
    else:
        self.flip_random_assignment()
else:
    logging.critical("naive_local_search timeout")
    return None
```

Local Search closing words

- Local Search is making a comeback in state-of-the-art solvers
- Algorithms like ProbSAT are almost competitive with... good solvers
- Further optimizations (that I did not implement)
 - Full restarts (reassign all variables)
 - Keep track of how much each variables contribute to conflicts in order to pick variables smarter

DPLL solvers

- Abbreviation of names: Davis-Putnam-Logemann-Loveland
- “Lazy Loop.py”
- Instead of assigning all literals at once, assign one and propagate everything you can.
- If we find a conflict before all variables are assigned, we can avoid assigning more variables!

Unit propagation


- If a clause is not satisfied and only has one unassigned literal, that literal must be satisfied
- Example:
 - $(\textcolor{red}{x1} \vee \neg \textcolor{red}{x2} \vee x3)$
 - $x3$ **must** be $\textcolor{green}{T}$ to satisfy the clause
- After unit propagation, more clauses may become unit!
- Recursively apply unit propagation until
 - Conflict \rightarrow UNSAT
 - All clauses satisfied \rightarrow SAT
 - Nothing can be propagated \rightarrow Assign some new literal

Implementation: DPLL.py

```
def DPLL(self, model0):
    sat, model1 = self.unit_propagation(model0)
    if sat is not None:
        return sat, model1
    # Pick an unassigned literal
    assigned_literals = list(map(abs, model1))
    unassigned_literals = [lit for lit in self.literals()
                           if lit not in assigned_literals]
    picked_literal = unassigned_literals[0]

    # Set it to True and recurse
    modelTrue0 = self.assign(model1, picked_literal)
    sat, modelTrue1 = self.DPLL(modelTrue0)
    if sat:
        return sat, modelTrue1

    # Set it to False and recurse
    modelFalse0 = self.assign(model1, -picked_literal)
    return self.DPLL(modelFalse0)
```



- Loops: search models at the leaves
- DPLL: descend slowly through the tree and check if the current model is already SAT/UNSAT

CDCL solvers


- Conflict Driven Clause Learning
- State of the art
- Like DPLL, but when you find a conflict you figure out what caused it and add a new clause to avoid it
 - Clause learning!
- Introducing shortcuts in the search tree
- Minisat is a CDCL solver

Decision level

- Decision level is the number of literals that have been arbitrarily assigned
- Notation:
 - 1 @ 3: variable 1 (x1) is assigned True at decision level 3
 - -5 @ 2: variable 5 (x5) is assigned False at decision level 2
- If we have a conflict at decision level 0:
 - We have assigned nothing, yet we have a conflict \Rightarrow return UNSAT
- Our CDCL implementation is not recursive (which the DPLL implementation is), so we need to track this.

Implication Graph

- We need to track **why** a particular literal was assigned and **when**
 - When: Decision level
 - Why: Implications
- Literal assignment can be:
 - **arbitrary** (once per decision level, “pick an unassigned variable”) or
 - **forced** (unit propagation)
- The implication graph tracks both **why** and **when**



Implication graph / CDCL example


- $(x_1 \vee x_2 \vee x_3) \wedge (x_2 \vee \neg x_4) \wedge (x_2 \vee \neg x_3 \vee x_4)$

Implication graph / CDCL example

- $(x_1 \vee x_2 \vee x_3)$
 $(x_2 \vee \neg x_4)$
 $(x_2 \vee \neg x_3 \vee x_4)$

Implication graph / CDCL example



- $(\textcolor{red}{x1} \vee x2 \vee x3)$
 $(x2 \vee \neg x4)$
 $(x2 \vee \neg x3 \vee x4)$



arbitrary assignment

Implication graph / CDCL example


- $(x_1 \vee x_2 \vee x_3)$
 $(x_2 \vee \neg x_4)$
 $(x_2 \vee \neg x_3 \vee x_4)$



arbitrary assignment

Implication graph / CDCL example


- $(x_1 \vee x_2 \vee x_3)$
 $(x_2 \vee \neg x_4)$
 $(x_2 \vee \neg x_3 \vee x_4)$



unit propagation

Implication graph / CDCL example


- $(x_1 \vee x_2 \vee x_3)$
 $(x_2 \vee \neg x_4)$
 $(x_2 \vee \neg x_3 \vee x_4)$



unit propagation

Implication graph / CDCL example


- $(x_1 \vee x_2 \vee x_3)$
 $(x_2 \vee \neg x_4)$
 $(x_2 \vee \neg x_3 \vee x_4)$



Conflict! Why?

Implication graph / CDCL example


- $(x_1 \vee x_2 \vee x_3)$
 $(x_2 \vee \neg x_4)$
 $(x_2 \vee \neg x_3 \vee x_4)$



Conflict! Why?

Implication graph / CDCL example

- $(x_1 \vee x_2 \vee x_3)$
 $(x_2 \vee \neg x_4)$
 $(x_2 \vee \neg x_3 \vee x_4)$




$-2 \wedge 3$ is the problem

New clause: $\neg(-2 \wedge 3) = 2 \vee -3$

Implication graph / CDCL example



- $(x_2 \vee \neg x_3)$
 $(\textcolor{red}{x}_1 \vee x_2 \vee x_3)$
 $(x_2 \vee \neg x_4)$
 $(x_2 \vee \neg x_3 \vee x_4)$



new clause in place!
roll back and reassign x_2

Implication graph / CDCL example

- $(x_2 \vee \neg x_3)$
 $(x_1 \vee x_2 \vee x_3)$
 $(x_2 \vee \neg x_4)$
 $(x_2 \vee \neg x_3 \vee x_4)$




SAT!

We did not need the learned clause this time.

The implication graph is needed for conflict analysis.

Implementation: CDCL.py

- The model is tracked by a custom ImplicationGraph data structure, which can be visualized because fun
- Essentially a list of [dstnode, [srcnode]]
- Unit propagation is essentially unchanged
- Conflict analysis was tricky, but straightforward once you know what to do
- Simple restart algorithm



Out of scope SAT things

- Watched Literals
 - Lazy pointer data structure, for avoiding searching the entire clause
 - Lazy: do nothing upon backjumping
- Variable State Independent Decaying Sum
 - Lazy data structure for choosing the next literal to assign
- SAT/UNSAT mode
 - SAT and UNSAT instances benefit from different techniques
- Any intelligent choice of next variable
- Real-world examples

SAT take-aways

- It's the simplest hard problem
 - Other hard problems can be formalized as SAT, and then solved efficiently
 - Or rather, in theory just as hard, but efficiently in practice most of the time
- State of the art solvers are crazy good
 - And just get faster and faster due to clever algorithms and heuristics
- Modern solvers are CDCL-based
- BMC was the “killer app” for SAT

SMT

Satisfiability Modulo Theories

SMT: Satisfiability Modulo Theories

- SAT with richer logics (Theories) on top, such as
- EUF: Equality and Uninterpreted Functions
 - $x_1 = x_2 \Rightarrow f(x_1) = f(x_2)$
 - Workhorse: Congruence Closure
- LIA: Linear Integer Arithmetic
 - $(x_1 \geq 0) \wedge (x_1 \leq x_2)$
 - Workhorse: Simplex
- Arrays
- Combining theories
- Quantifiers

Example with EUF and LIA

- $(x_1 \geq 0) \wedge (x_1 < 1) \wedge ((x_2 = x_1) \vee (x_2 = 0)) \wedge$
 $(f(x_1) = f(x_2) \Rightarrow \neg (g(x_1) = g(x_2)))$
- $x_1 = 0$
- $x_2 = 0$
- $x_1 = x_2$
- $f(x_1) = f(x_2)$
- $g(x_1) = g(x_2)$
- $\neg (g(x_1) = g(x_2))$
- $\rightarrow \text{UNSAT}$

Lingua Franca: smtlib

```
(set-logic QF_LRA) ; Reals
(declare-const a Real)
(declare-const b Real)
(declare-const c Real)
(declare-const d Real)
(assert (= 1 (+ (* 2 a) c))) ;  $2a + c = 1$ 
(assert (= 0 (+ (* 2 b) d))) ;  $2b + d = 0$ 
(assert (= 0 (+ (* 2 c) a))) ;  $a + 2c = 0$ 
(assert (= 1 (+ (* 2 d) b))) ;  $b + 2d = 1$ 
(check-sat)
(get-model)
```

Python interface

```
1 %reset -f
2 from z3 import *
3 M = [[2, 1],
4       [1, 2]]
5 # Declare four variables over the reals to represent the inverse matrix M^-1
6 a, b, c, d = Reals("a b c d")
7 Minv = [[a, b],
8          [c, d]]
9 # Formulate four equations characterising the inverse, use the function solve
10 eq1 = M[0][0]*Minv[0][0] + M[0][1]*Minv[1][0] == 1
11 eq2 = M[0][0]*Minv[0][1] + M[0][1]*Minv[1][1] == 0
12 eq3 = M[1][0]*Minv[0][0] + M[1][1]*Minv[1][0] == 0
13 eq4 = M[1][0]*Minv[0][1] + M[1][1]*Minv[1][1] == 1
14 solve(And(eq1, eq2, eq3, eq4))

[c = -1/3, a = 2/3, d = 2/3, b = -1/3]
```

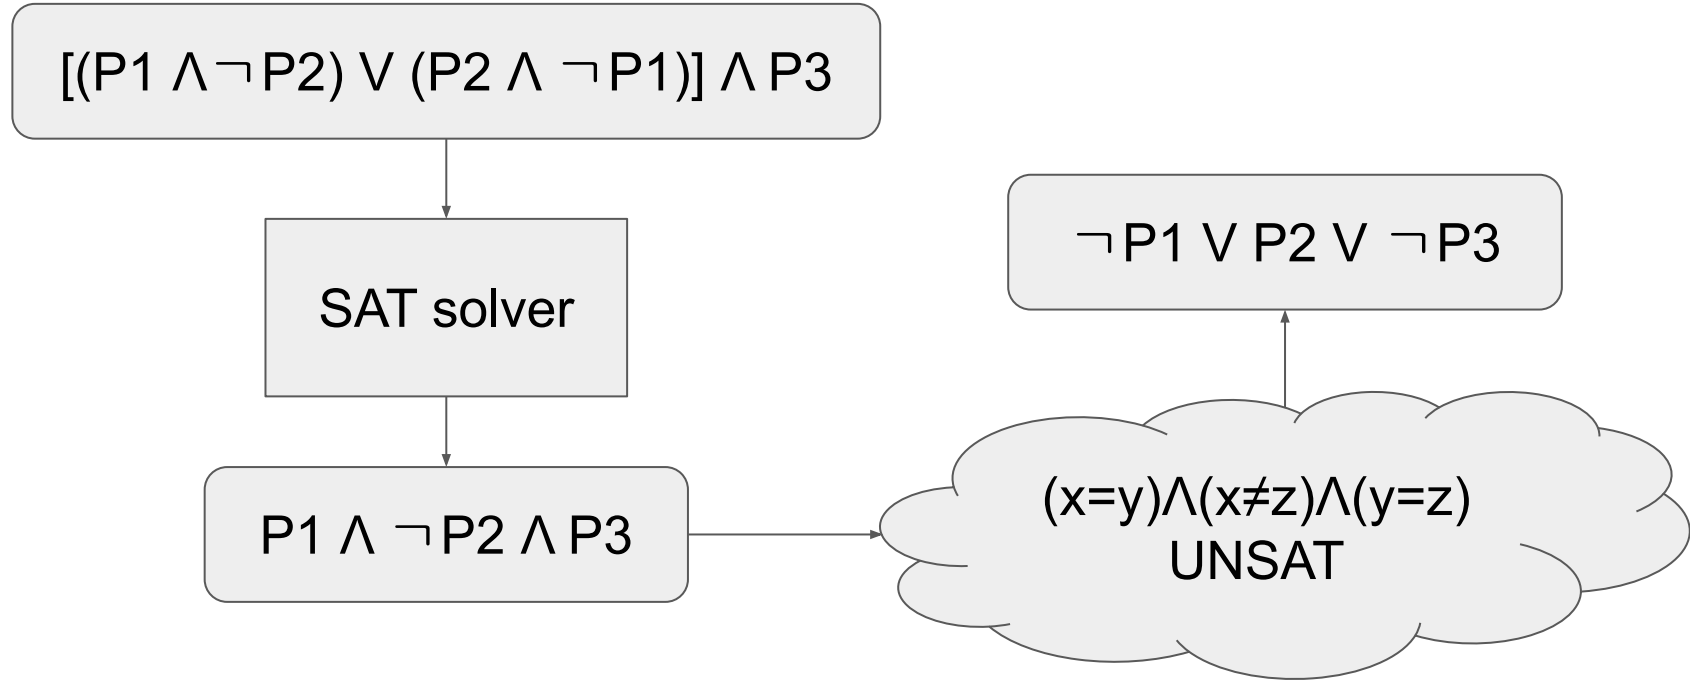
Lazy SMT

- Principle: *deciding satisfiability of a formula can be reduced to deciding the theory satisfiability of conjunctions of constraints.*
- Example: $[(x=y) \wedge (x \neq z) \vee (x=z) \wedge (x \neq y)] \wedge (y=z)$
- Assign $P1 := x=y$ $P2 := x=z$ $P3 := y=z$
- Formula: $[(P1 \wedge \neg P2) \vee (P2 \wedge \neg P1)] \wedge P3$
- DNF: $(P1 \wedge \neg P2 \wedge P3) \vee (\neg P1 \wedge P2 \wedge P3)$
 - Check each conjunction separately (theory satisfiability)
 - $(P1 \wedge \neg P2 \wedge P3)$: $(x=y) \wedge (x \neq z) \wedge (y=z)$
 - $(\neg P1 \wedge P2 \wedge P3)$: $(x \neq y) \wedge (x=z) \wedge (y=z)$
- All conjuncts are UNSAT, so the expression is UNSAT.

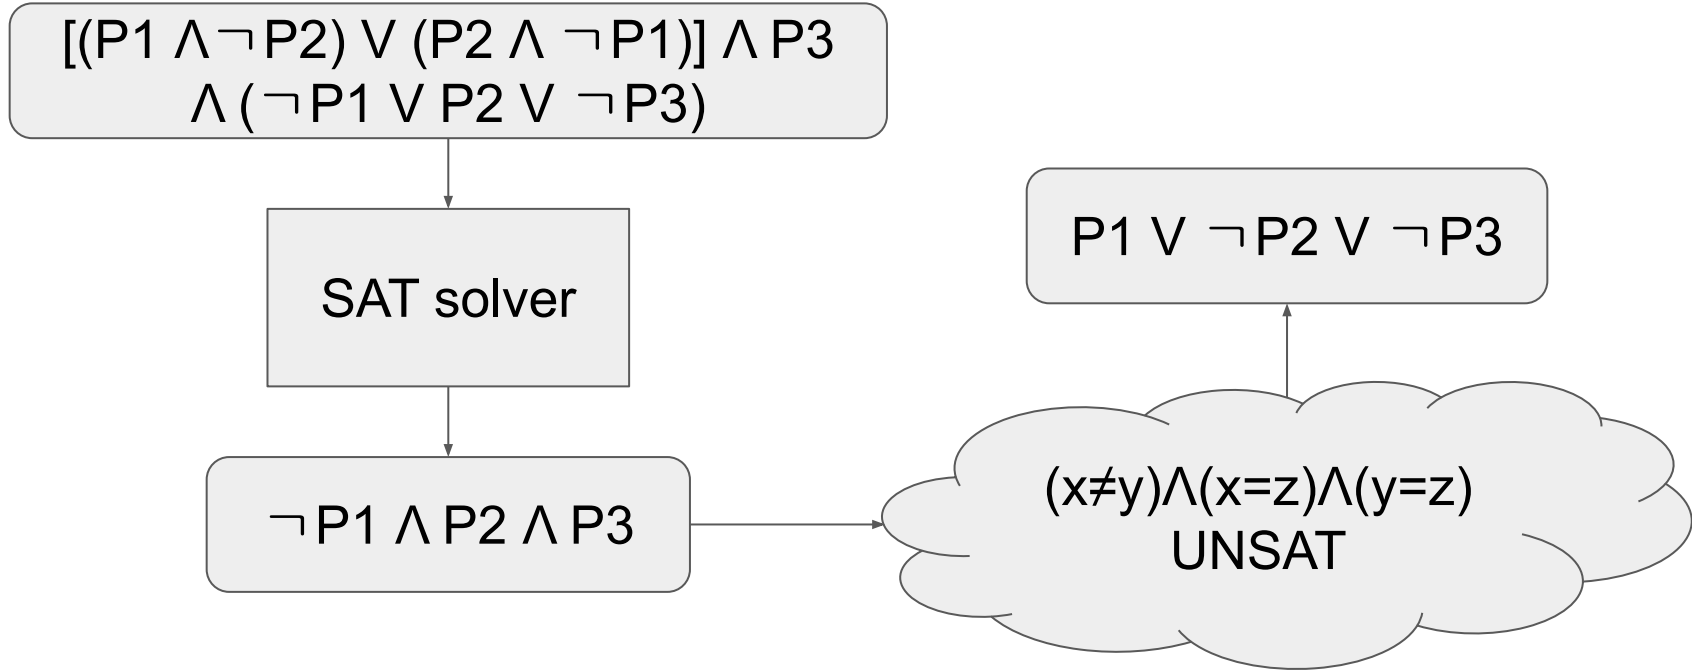
Lazy SMT

- Principle: *deciding satisfiability of a formula can be reduced to deciding the theory satisfiability of conjunctions of constraints.*
- **Converting to DNF is too expensive.** Can we avoid it?
- Lazy SMT: use a SAT solver to enumerate conjuncts!

Lazy SMT



Lazy SMT



Lazy SMT

$$[(P1 \wedge \neg P2) \vee (P2 \wedge \neg P1)] \wedge P3$$
$$\wedge (\neg P1 \vee P2 \vee \neg P3)$$
$$\wedge (P1 \vee \neg P2 \vee \neg P3)$$

SAT solver

UNSAT

CDCL(T)

- Lazy SMT: modular SAT and T solvers, “offline”
- CDCL(T): “online” approach, tighter integration between solvers
 - What should you give T solvers?
 - Can T solvers prune the search space?
 - Propagation?
 - ...
- Out of scope for today

The eager approach, “bit blasting”

- Convert as much as possible into pure SAT
- Usually the best and fastest way, especially for bit vector arithmetic

Implementation project idea

- Choose the simplest logic
 - EUF
- Implement a solver for that logic
 - Congruence Closure
- Use our SAT solver to find conjunctions
- \Rightarrow Lazy SMT

EUF: Equality and Uninterpreted Functions

EUF syntax

(Formula) $\varphi ::= Atom \mid \varphi \wedge \varphi \mid \varphi \vee \varphi \mid \neg \varphi$

(Atom) $Atom ::= Term = Term$

(Term) $Term ::= Var \mid Const \mid F(Term)$

Reflexivity: $\frac{}{E = E}$ Transitive: $\frac{E_1 = E_2 \quad E_2 = E_3}{E_1 = E_3}$

Symmetry: $\frac{E_2 = E_1}{E_1 = E_2}$ Congruence: $\frac{E_1 = E_2}{f(E_1) = f(E_2)}$


EUUF examples

- $x_1 = x_2$
 - SAT
- $(x_1 = x_2) \wedge (F(x_1) \neq F(x_2))$
 - UNSAT
- $(x_1 = x_2) \wedge (x_2 = x_3) \wedge (G(x_1) \neq G(x_3))$
 - UNSAT
- $(x_1 = x_2 \vee x_3 = x_4) \wedge \neg (F(x_1) = F(x_2) \vee G(x_3) = G(x_4))$
 - UNSAT

Congruence Closure

Congruence Closure


- $x1 = x2$



New equivalence class

Congruence Closure


- $x1 = x2$
- $x2 = x3$



New equivalence class

Congruence Closure

- $x1 = x2$
- $x2 = x3$




If any classes share a term, merge them

Congruence Closure


- $x1 = x2$
- $x2 = x3$
- $f(x1) = g(x1)$

New equivalence class



Congruence Closure


- $x_1 = x_2$
- $x_2 = x_3$
- $f(x_1) = g(x_1)$
- $f(x_2) = g(x_4)$



New equivalence class

Congruence Closure

- $x_1 = x_2$
- $x_2 = x_3$
- $f(x_1) = g(x_1)$
- $f(x_2) = g(x_4)$




Congruence:

$x_1 = x_2$ so $f(x_1) = f(x_2)$

Congruence Closure


- $x_1 = x_2$
- $x_2 = x_3$
- $f(x_1) = g(x_1)$
- $f(x_2) = g(x_4)$
- $f(x_1) \neq x_1$



SAT, not in same class

Congruence Closure

- $x1 = x2$
- $x2 = x3$
- $f(x1) = g(x1)$
- $f(x2) = g(x4)$
- $f(x1) \neq x1$
- $f(x1) \neq g(x4)$

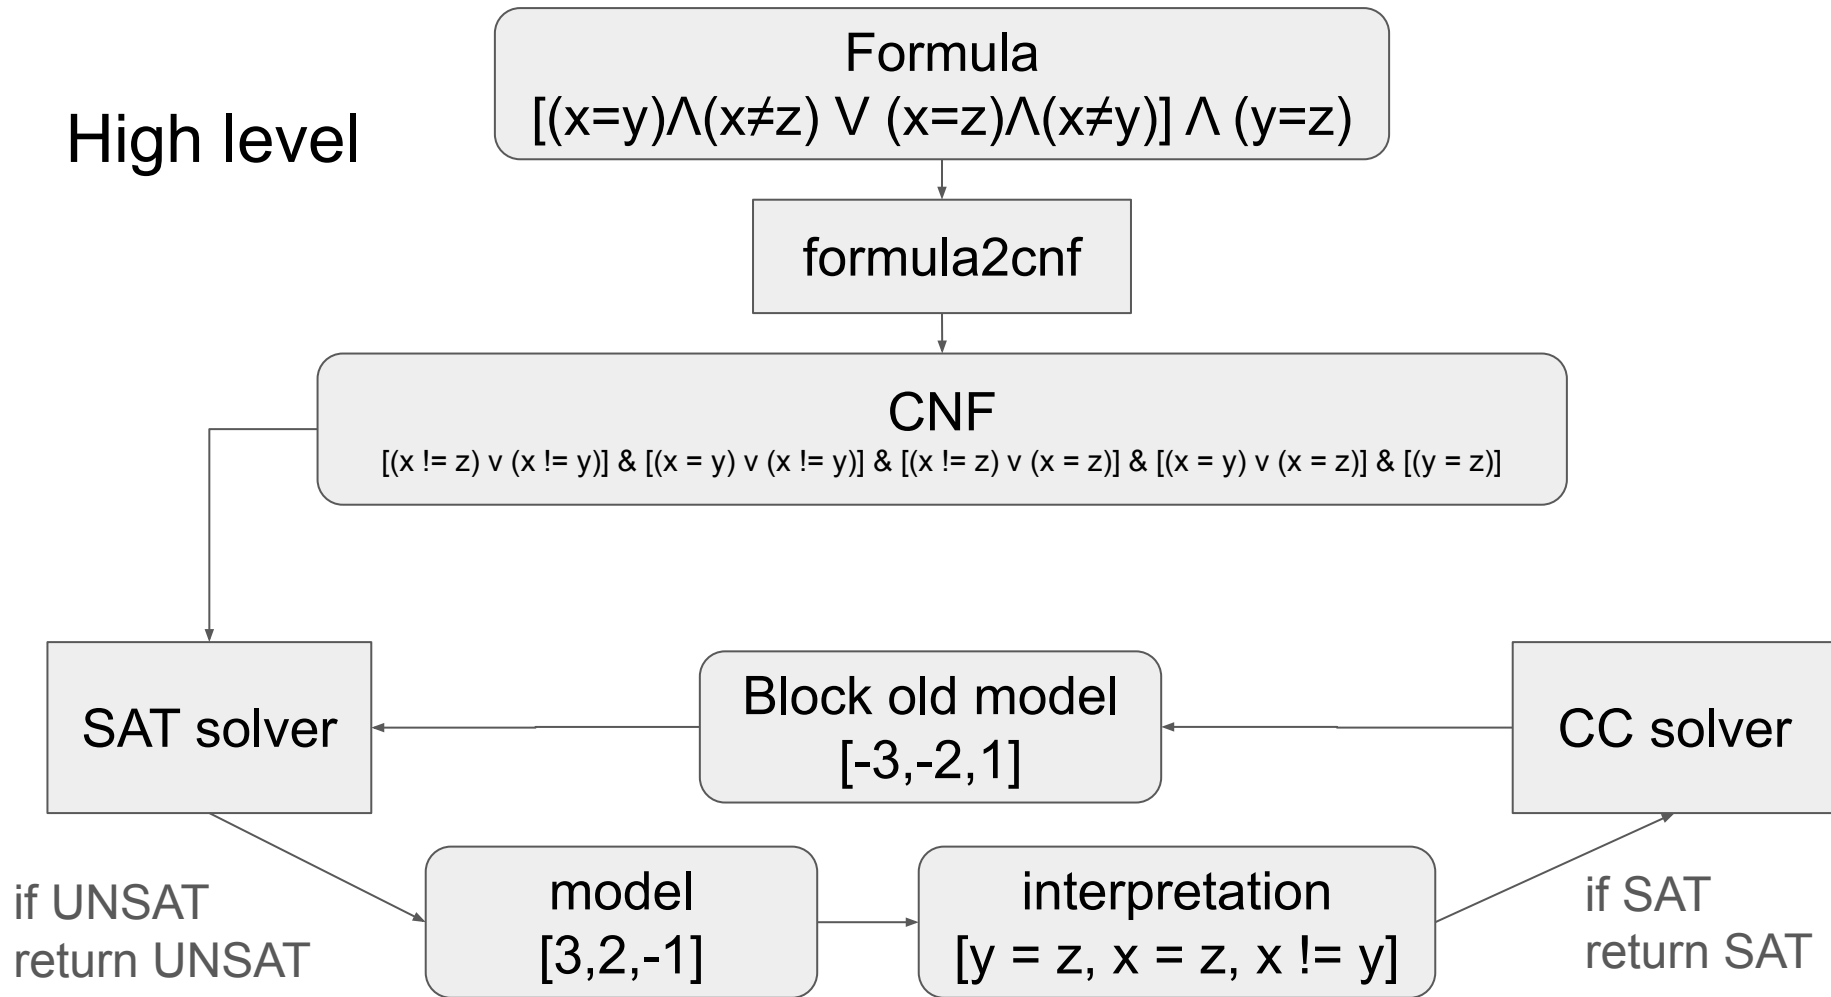


UNSAT, $f(x1) = g(x4)$

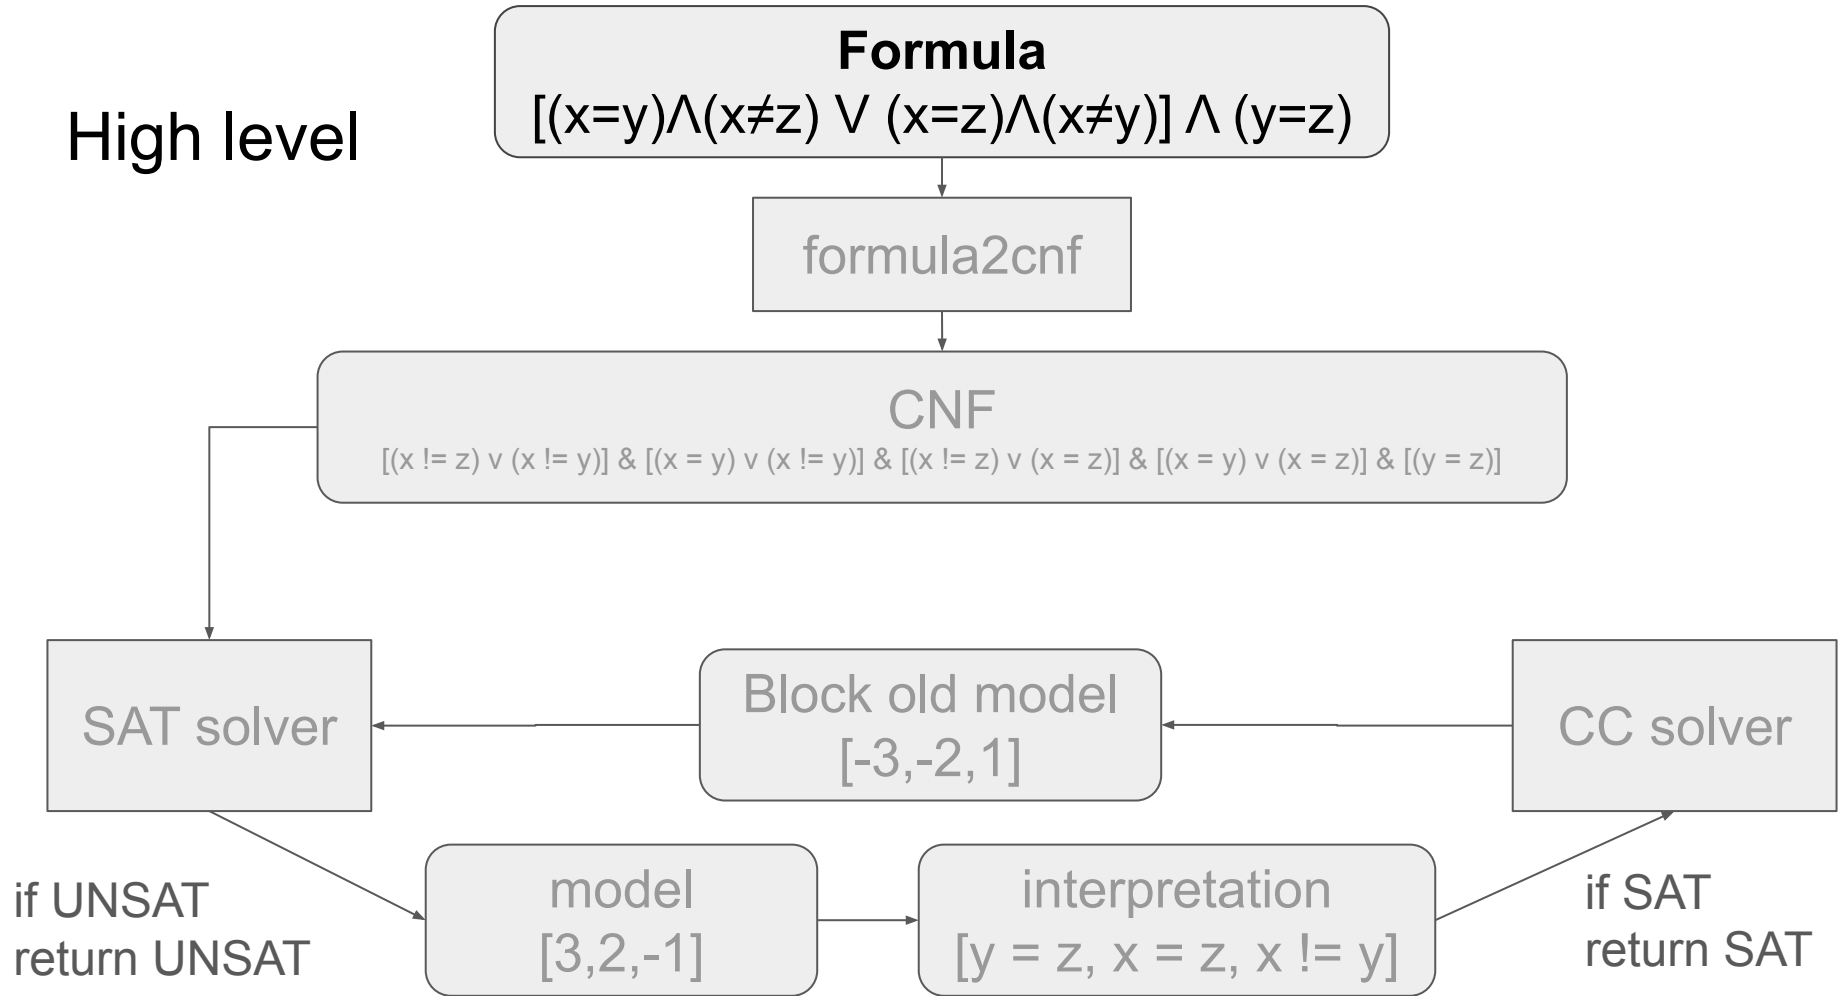
Implementation: CongruenceClosure.py

- This was intended to be a small thing
- 1000+ lines of python later, that turned out not to be the case

High level



High level



Formula

EUf syntax

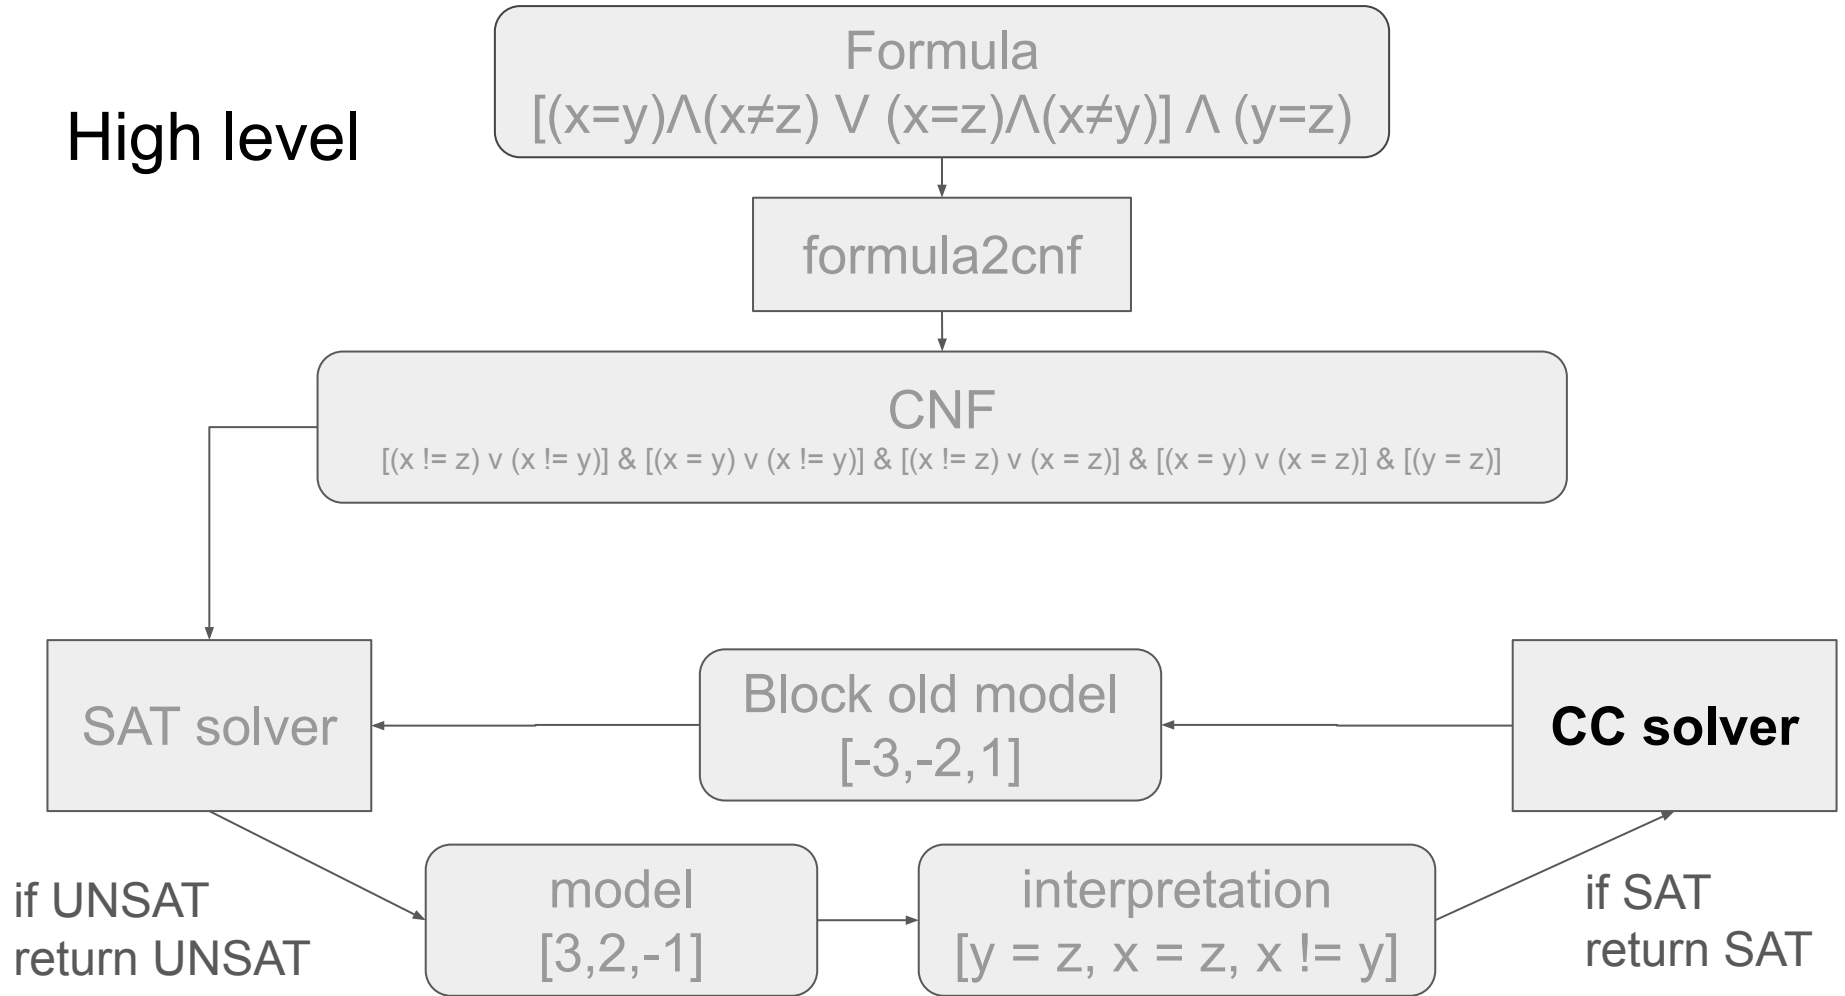
(Formula) $\varphi ::= Atom \mid \varphi \wedge \varphi \mid \varphi \vee \varphi \mid \neg \varphi$

(Atom) $Atom ::= Term = Term$

(Term) $Term ::= Var \mid Const \mid F(Term)$

- General formulas for EUF

High level



SimpleCongruenceClosure

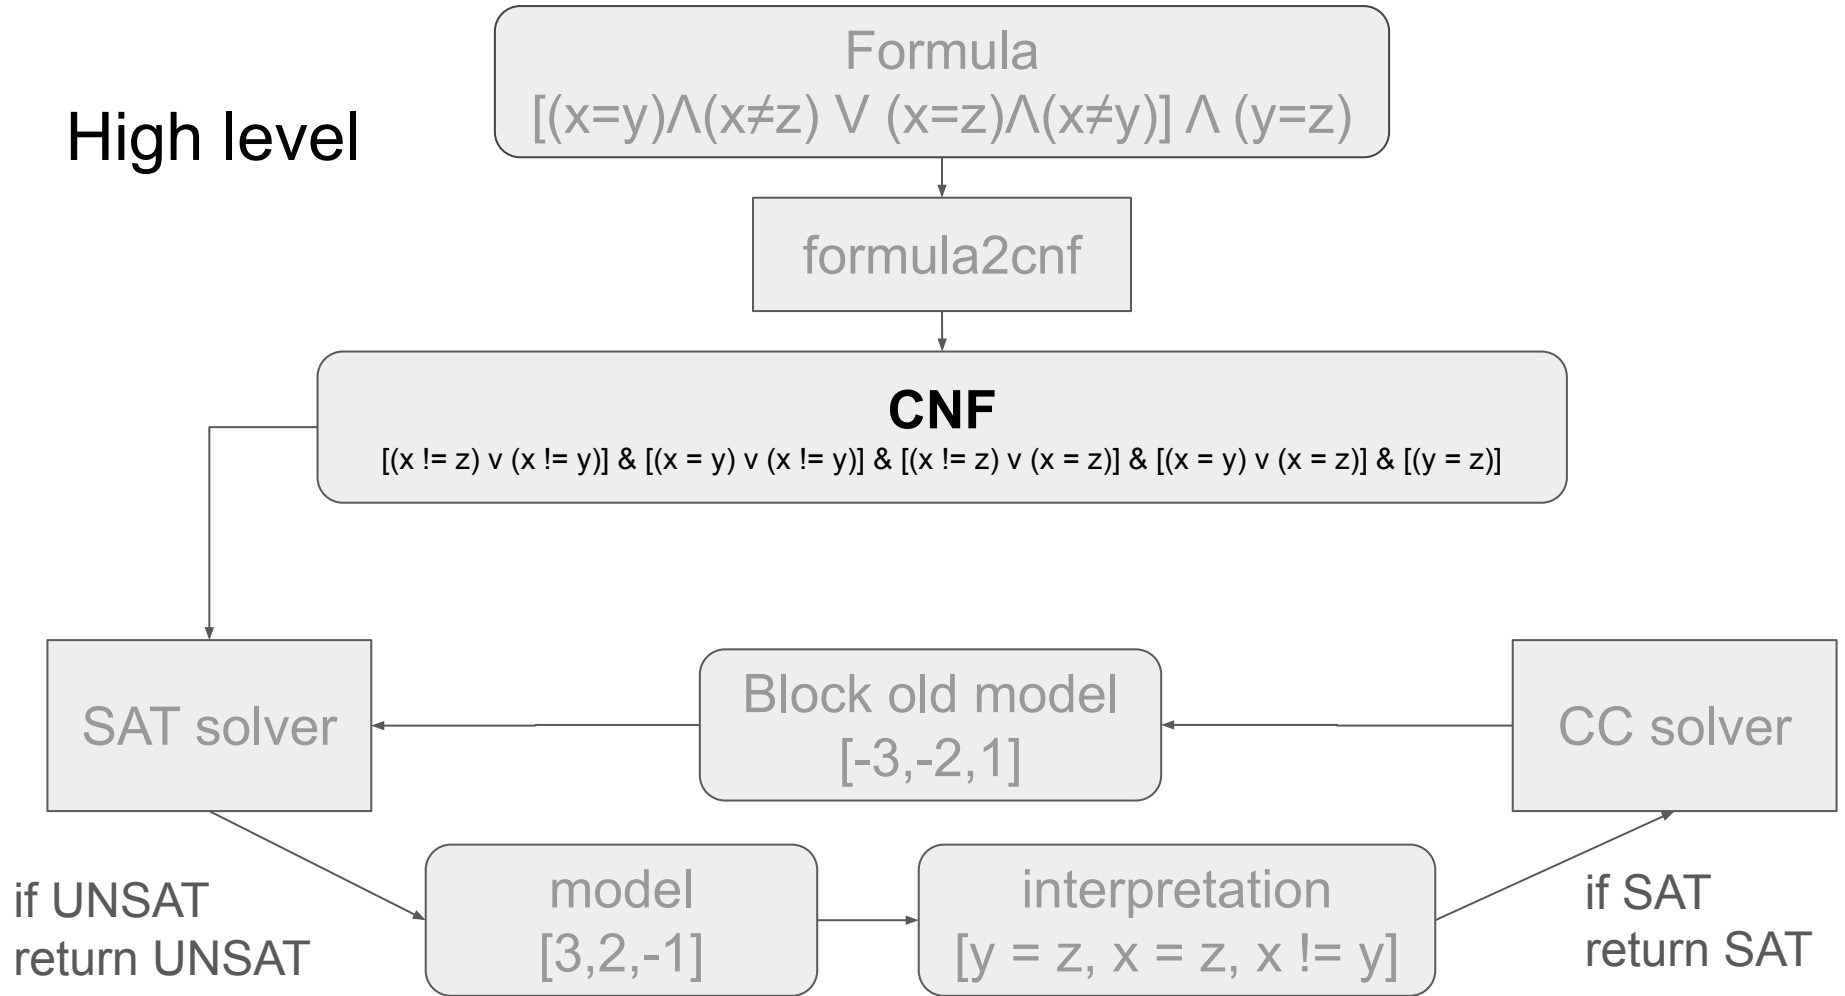
- Input: Conjunctions of
Term = Term | Term \neq Term
 - Term is
Var | F(Term)
- Output: SAT/UNSAT

```

<CC solver>
All (dis)equalities:
    F(x1) != F(x2)
CC_solv    x1 = x2
All (di    Unique terms:
    x !=
    x =    x1
    y =
Unique     x2
    y     F(x1)
    z
    x     F(x2)
Disequa   Disequalities
    x !=
Equival    F(x1) != F(x2)
    y
    z     Equivalence class 0:
    x
</CC so    x1
Checkin    x2
testing
    aga</CC solver>
testing
    agaUNSAT
Bottom!    Should be UNSAT

```

High level

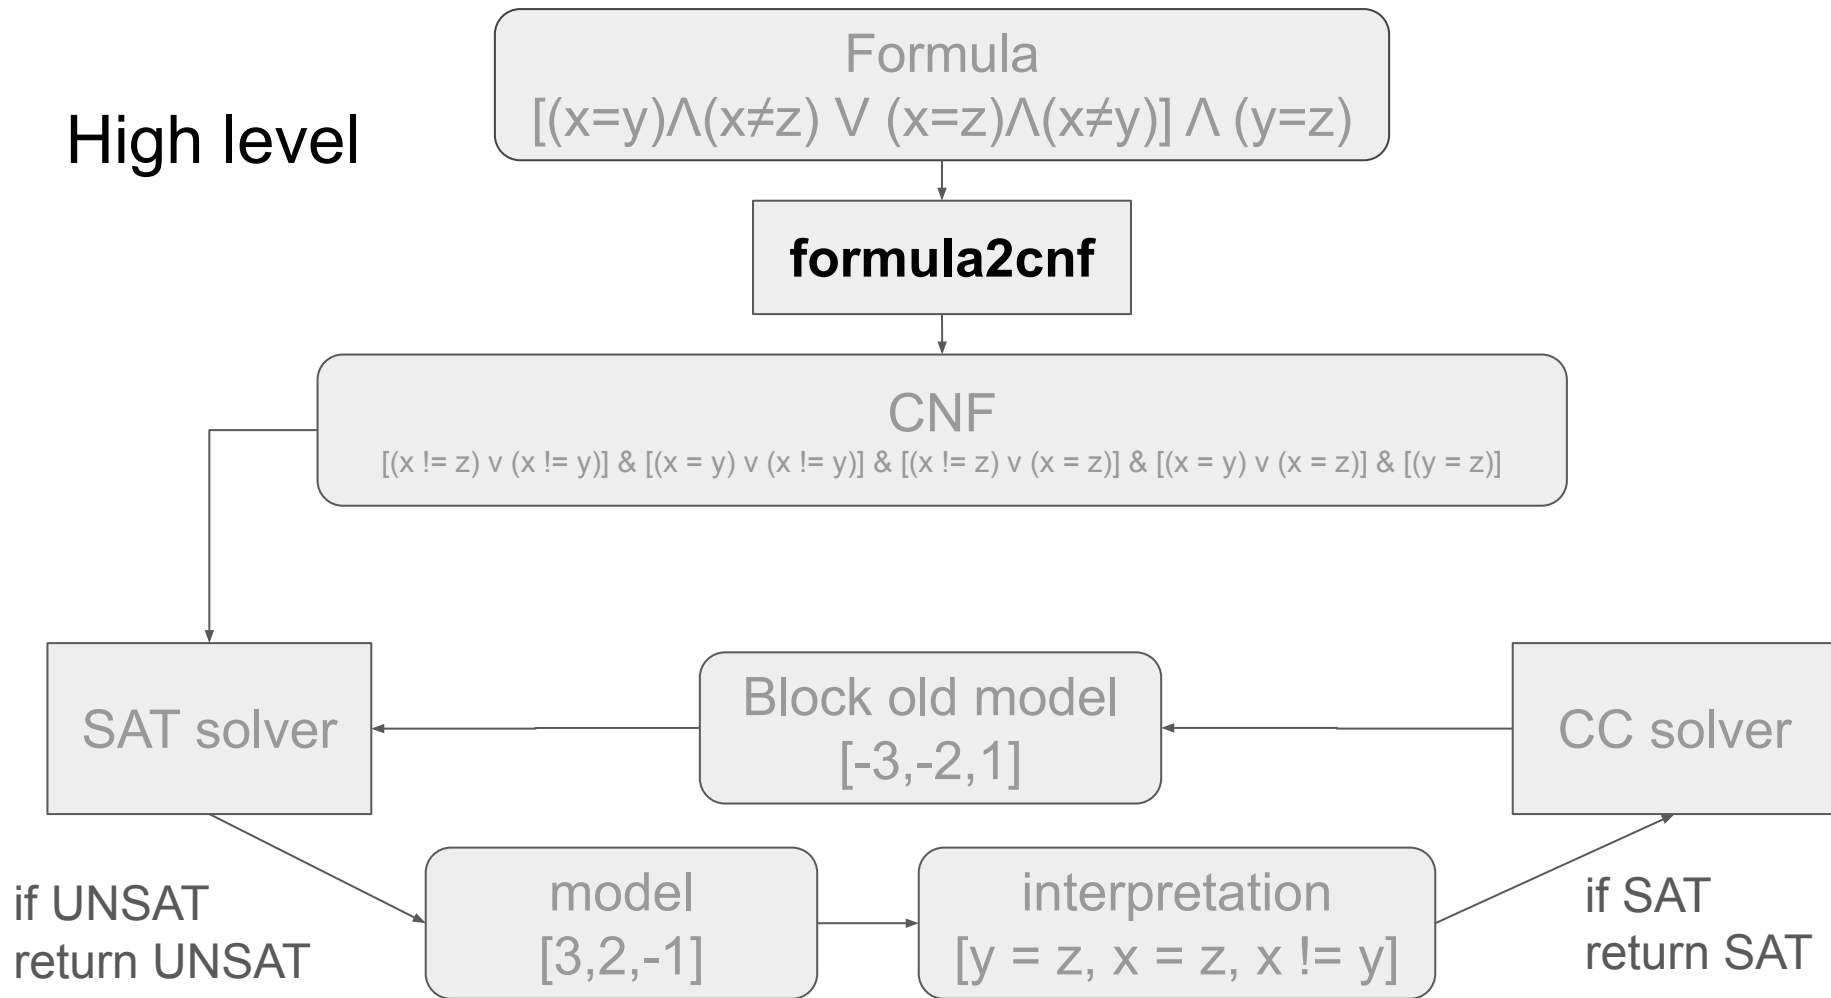


CNF

- A class with a list of conjuncts
- Literals are also on the form
Term=Term | Term≠Term
- and(cnf1, cnf2): trivial
- not(cnf): DeMorgan
- or(cnf1, cnf2): distribution

```
c1: [(a = a) v (b = b)] & [(c = c) v (d = d)]  
c2: [(e = e) v (f = f)] & [(g = g) v (h = h)]  
c1 or c2: [(e = e) v (f = f) v (a = a) v (b = b)] & [(g = g) v (h = h) v (a = a) v (b  
= b)] & [(e = e) v (f = f) v (c = c) v (d = d)] & [(g = g) v (h = h) v (c = c) v (d =  
d)]
```

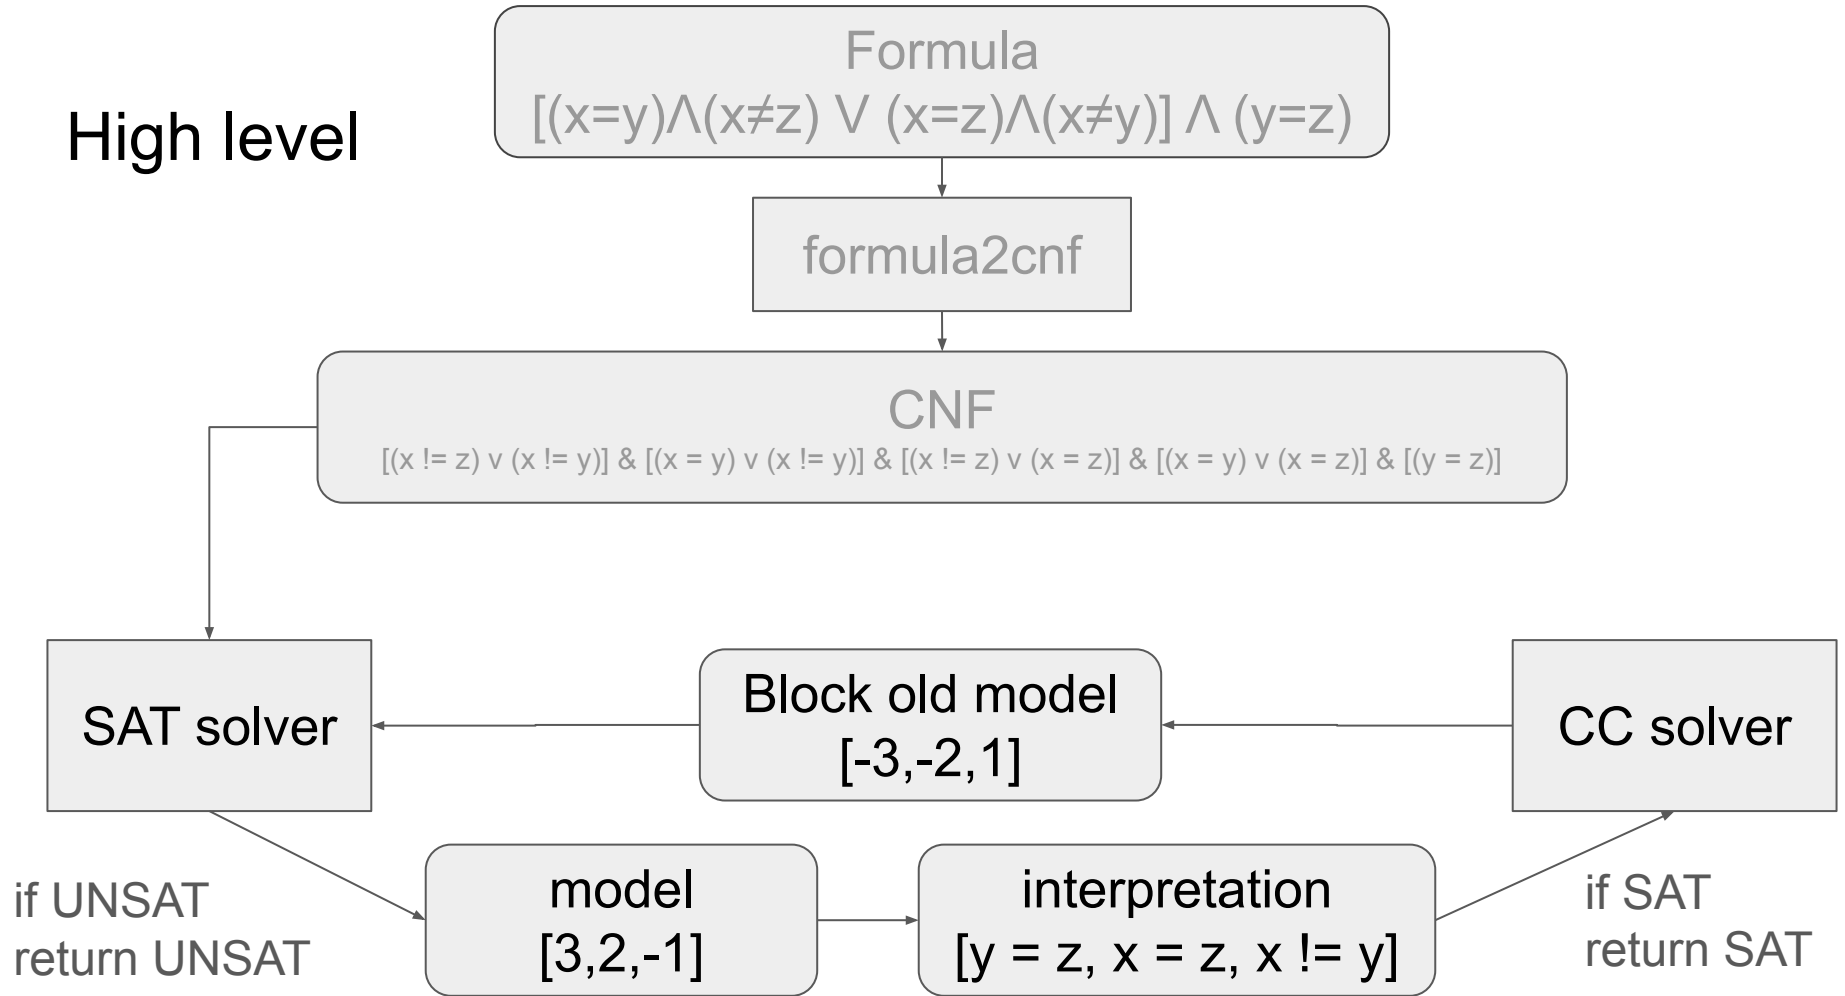
High level



formula2cnf

- Naive conversion to cnf is simple
 - Push all negations down to literals
 - Exponential blow-up
- Proper way: Tseitin algorithm
 - Introduce “extra” variables
 - Linear size
 - Satisfiability-preserving
- In the interest of time: Naive way

High level



Solver loop

```
# The SAT model is invalid.  
# Negate it and add it as a clause to the SAT solver.  
neg_model_clause = []  
for literal in sat_model:  
    neg_model_clause.append(literal * -1)  
if verbose:  
    print(f"neg_model_clause: {neg_model_clause}")  
SAT_solver.add_clause(neg_model_clause)
```

A word about ordering

- I did not want to deal with symmetry etc: $a=b \Leftrightarrow b=a$
- Impose a total ordering for each class,
enforced when the class is instantiated
- `Atom(a,b)` and `Atom(b,a)` will create the same thing: $a=b$

SMT things not covered

- A lot.
- Main SMT solvers: z3 and cvc5
 - Both have python wrappers that easy to use
- Quantifiers!
 - Pose a huge challenge for SMT solvers
- Other Theories
- Real-world examples